

EDUCATION

Purdue University

Ph.D. in Computer Science

Advisor: Dr. Ruqi Zhang

IN, USA

Aug. 2025 – Present

Tianjin University

B.S. in Data Science and Big Data Technology

School of Mathematics · Advisors: Dr. Bing Cao and Dr. Qinghua Hu

Tianjin, China

Aug. 2021 – Jun. 2025

RESEARCH INTERESTS

I study how large language and vision-language models can improve from limited, noisy, or imperfect supervision. My work develops reliable and efficient post-training algorithms for reasoning and safety, with a focus on reinforcement learning, on-policy distillation, and self-correction.

PUBLICATIONS

* Equal contribution

Preprints and Technical Reports

[P1] Does On-Policy Distillation Really Distill? From Noisy Teacher to Self-Improvement

Yi Ding, Ruqi Zhang

Preprint, 2026 [PDF] [Project]

#1 on Hugging Face Daily Papers (Sep. 1, 2026)

[P2] Addressing Performance Saturation for LLM RL via Precise Entropy Curve Control

Bolian Li, Yifan Wang, Yi Ding, Anamika Lochab, Ananth Grama, Ruqi Zhang

Preprint, 2026 [PDF] [Project]

[P3] SafeWork-R1: Coevolving Safety and Intelligence under the AI-45° Law

Shanghai AI Lab: Yicheng Bao, ..., Yi Ding, ... (100+ authors)

Technical report, 2025 [PDF]

Conference Publications

[C1] Learning Self-Correction in Vision–Language Models via Rollout Augmentation

Yi Ding, Ziliang Qiu, Bolian Li, Ruqi Zhang

ICML 2026 · International Conference on Machine Learning [PDF] [Project]

[C2] Position: Modular Safety Guardrails Are Necessary for Foundation-Model-Enabled Robots in the Real World

Joonkyung Kim, Wenxi Chen, Davood Soleymanzadeh, Yi Ding, Xiangbo Gao, Zhengzhong Tu, Ruqi Zhang, Fan Fei, Sushant Veer, Yiwei Lyu, Minghui Zheng, Yan Gu

ICML 2026 · International Conference on Machine Learning [PDF]

[C3] Rethinking Bottlenecks in Safety Fine-Tuning of Vision Language Models

Yi Ding*, Lijun Li*, Bing Cao, Jing Shao

ICLR 2026 · International Conference on Learning Representations [PDF] [Project]

PUBLICATIONS (CONTINUED)

[C4]

Sherlock: Self-Correcting Reasoning in Vision-Language Models

Yi Ding, Ruqi Zhang

NeurIPS 2025 · Advances in Neural Information Processing Systems [PDF] [Project]

[C5]

Visual Contextual Attack: Jailbreaking MLLMs with Image-Driven Context Injection

Ziqi Miao*, Yi Ding*, Lijun Li, Jing Shao

EMNLP 2025 · Empirical Methods in Natural Language Processing, Main Conference [PDF]

[C6]

ETA: Evaluating Then Aligning Safety of Vision Language Models at Inference Time

Yi Ding, Bolian Li, Ruqi Zhang

ICLR 2025 · International Conference on Learning Representations [PDF] [Project]

[C7]

Test-Time Dynamic Image Fusion

Bing Cao, Yinan Xia*, Yi Ding*, Changqing Zhang, Qinghua Hu

NeurIPS 2024 · Advances in Neural Information Processing Systems [PDF]

[C8]

Predictive Dynamic Fusion

Bing Cao, Yinan Xia*, Yi Ding*, Changqing Zhang, Qinghua Hu

ICML 2024 · International Conference on Machine Learning [PDF]

RESEARCH EXPERIENCE

RZ-Lab, Purdue University

Research Assistant

Advisor: Dr. Ruqi Zhang

IN, USA

May 2024 – May 2025

Open Trust Lab, Shanghai AI Laboratory

Research Scientist

Advisor: Dr. Lijun Li

Beijing, China

Dec. 2024 – Mar. 2025

MLDM Lab, Tianjin University

Research Assistant

Advisors: Dr. Bing Cao and Dr. Qinghua Hu

Tianjin, China

Sep. 2023 – Dec. 2024

AWARDS AND RESEARCH SUPPORT

NSF ACCESS Discover Project Award

Project: LLM Post-Training

2025–2026

ACADEMIC SERVICE

Reviewer: ICLR 2025–2027; ICML 2026; NeurIPS 2025, 2026; CIKM 2026; ACL Rolling Review (May 2025).

SKILLS AND LANGUAGES

Technical: Python, MATLAB, $\text{\LaTeX}$; mathematical and statistical analysis.

Languages: Mandarin Chinese (native); English (TOEFL: 102).